

Getting good performance from your application

Tuning techniques for serial programs on
cache-based computer systems

Application Tuning


Selected Topics

Application Tuning


- ❑ Selected Topics:
 - ❑ 32- vs 64-bit
 - ❑ binary data portability
 - ❑ floating point numbers and IEEE 754
 - ❑ compiler options
 - ❑ a case study
- ❑ Summary

32-bit vs 64-bit issues

- ❑ 64-bit operating systems
- ❑ Implication: The address space of a single application can be larger than 4 GB



32-bit vs 64-bit issues



- ❑ **Addresses \neq Data**
- ❑ An 'n'-byte data type fills always n bytes in memory (byte addressable)
- ❑ I.e. the next element is n bytes further in memory
- ❑ This increment is not related to the size of the addresses (32-bit or 64-bit)

32-bit vs 64-bit issues

<u><i>C data type</i></u>	<u><i>ILP32</i></u> <i>(bits)</i>	<u><i>LP64</i></u> <i>(bits)</i>
<i>char</i>	8	<i>same</i>
<i>short</i>	16	<i>same</i>
<i>int</i>	32	<i>same</i>
<i>long</i>	32	64
<i>long long</i>	64	<i>same</i>
<i>pointer</i>	32	64
<i>enum</i>	32	<i>same</i>
<i>float</i>	32	<i>same</i>
<i>double</i>	64	<i>same</i>
<i>long double</i>	128	<i>same</i>

UNIX and Linux support LP64; Windows 64-bit uses LLP64, where long stays 32 bits

(p)ldd and LD_LIBRARY_PATH

How to check which shared-libraries are loaded?

- ❑ Static check: use the ldd command

- ❑ `$ ldd executable`

- ❑ Dynamic check: use pldd on the PID

- ❑ `$ pldd pid`

- ❑ Solaris only

- ❑ there are scripts available for Linux as well

- ❑ we have installed pldd on the DTU HPC cluster

(p)ldd and LD_LIBRARY_PATH

- ❑ How to change the search path for dynamic libraries?
 - ❑ Use LD_LIBRARY_PATH – but use it with care!
- ❑ Best practice:
 - ❑ Compile the path into your application:
 - ❑ GCC: `-Wl,-rpath <path_to_lib>`
 - ❑ `ld.so` will then use this path
 - ❑ Avoid LD_LIBRARY_PATH in your shell environment – use a wrapper script for the application
- ❑ Check out this [blog note \(www.hpc.dtu.dk/?page_id=1180\)](http://www.hpc.dtu.dk/?page_id=1180), too!

Binary data storage

- ❑ Storing your data in binary format
- ❑ Advantages:
 - ❑ compact
 - ❑ fast
 - ❑ no loss of precision
- ❑ Drawbacks:
 - ❑ not “human readable”
 - ❑ data analysis more complicated
 - ❑ and ...

Binary data storage

- ❑ Example: integer 0x12345678 (hexadecimal)

```
value = 0x12345678;           // 305419896
printf("%d\n", value);
fwrite(&value, sizeof(value), 1, fptr);
```

- ❑ Write it ...

- ❑ ... on i386:

```
305419896
Architecture: i386
Value written to endian_i386.dat.
```

- ❑ ... on SPARC:

```
305419896
Architecture: sparc
Value written to endian_sparc.dat.
```

Binary data storage

❑ Read it:

```
fread(&value, sizeof(value), 1, fptr);  
printf("%d\n", value);
```

❑ on i386 data from i386:

```
Architecture: i386
```

```
Read from endian_i386.dat: 305419896
```

❑ on i386 data from SPARC:

```
Architecture: i386
```

```
Read from endian_sparc.dat: 2018915346
```




Little Endian vs Big Endian

- ❑ The order in which the bits are interpreted has not been standardized!
- ❑ Two 'popular' formats in use
 - ❑ Big Endian – SPARC, PowerPC, ...
 - ❑ Little Endian – Intel x86, AMD64, ...
- ❑ This is an issue when using the same binary data file on both platforms ...

Little Endian vs Big Endian

- Example: integer 0x12345678 (hexadecimal)



- Check with 'od' command:

```
$ od -x endian_sparc.dat
00000000 1234 5678
00000004

$ od -x endian_i386.dat
00000000 7856 3412
00000004
```

Little Endian vs Big Endian

- ❑ This is something you should be aware of when working with binary data!
- ❑ Tools:
 - ❑ Sun Fortran: -xfilebyteorder option
 - ❑ Portland Fortran compiler
 - ❑ swab() subroutine (low level)
 - ❑ ...

Floating point numbers & IEEE 754

Lesser known side effects of IEEE 754:

- ❑ Will this code run or fail?

```
#include <stdio.h>
#include <math.h>

int
main(int argc, char *argv[]) {

    double x;

    for(int i = 0; i < 10; i++) {
        x = sqrt(5.0 - i);
        printf("%lf\n", x);
    }
}
```

Floating point numbers & IEEE 754

Lesser known side effects of IEEE 754:

❑ What do you prefer?

```
$ gcc -otrapex trapex.c -lm
$ ./trapex
2.236068
2.000000
1.732051
1.414214
1.000000
0.000000
-nan
-nan
-nan
-nan
$
```

IEEE 754 compliant!

```
$ suncc -ftrap=common -o trapex
$ ./trapex
2.236068
2.000000
1.732051
1.414214
1.000000
0.000000
Floating point exception
$
```

not(!)
IEEE 754 compliant!

Floating point numbers & IEEE 754

Lesser known side effects of IEEE 754:

- ❑ The IEEE 754 standard doesn't “allow” traps on floating point exceptions, like invalid arguments, division by zero, over- and underflows
- ❑ Some compilers provide options to change that.
 - ❑ Studio: -ftrap=<exception_list>, e.g. common
 - ❑ Intel: -fp-trap=<exception_list>, e.g. common
- ❑ However: GCC has no such option, needs to be implemented by the programmer via library calls (see ‘man fenv’ – and next slide)

Floating point numbers & IEEE 754

Enabling floating point traps with GCC

```
#include <stdio.h>
#include <math.h>

#define __USE_GNU
#include "fenv.h"

void enable_ftraps(void) {
    feenableexcept (FE_DIVBYZERO);
    feenableexcept (FE_INVALID);
    feenableexcept (FE_OVERFLOW); }

int main(int argc, char *argv[]) {
    double x;
    enable_ftraps();
    for(int i = 0; i < 10; i++) {
        x = sqrt(5.0 - i);
        printf("%lf\n", x);
    }
}
```

Floating point numbers & IEEE 754

Now it works with GCC ...

```
$ gcc -o trapex2 trapex2.c -lm
$ ./trapex2
2.236068
2.000000
1.732051
1.414214
1.000000
0.000000
Floating point exception
$
```

not(!)
IEEE 754 compliant!

Floating point numbers & IEEE 754

- ❑ Compilers do not re-arrange your arithmetic expressions – unless you ask for it!
- ❑ This is part of the optimization flags (and varies from compiler to compiler!)
- ❑ Example on the next slides shows this for a division inside a for-loop
 - ❑ this is done with an old Studio compiler, but illustrates nicely what happens
 - ❑ the optimization option is ‘-fsimple=[0|1|2]’, from none (0) to aggressive arithmetic optimizations (2).

Floating point numbers & IEEE 754

Effects of -fsimple:

- ❑ compiled with -fast -xrestrict -fsimple=0:

```
1. void
2. divvec(int n, double div, double *a, double *b) {
3.
4.     int i;
5.
6.     for(i = 0; i < n; i++)
7.         b[i] = a[i]/div;
8. }
```

Source loop below has tag L1
L-1 scheduled with steady-state cycle count = 17
L-1 has 1 loads, 1 stores, 2 prefetches, 0 FPadds,
0 FPMuls, and 1 FPdivs per iteration

UltraSPARC

Floating point numbers & IEEE 754

Effects of -fsimple:

- ❑ compiled with -fast -xrestrict -fsimple=2:

```
1. void
2. divvec(int n, double div, double *a, double *b) {
3.
4.     int i;
5.
6.     for(i = 0; i < n; i++)
7.         b[i] = a[i]/div;
8. }
```

Source loop below has tag L1
L-1 scheduled with steady-state cycle count = 1
L-1 unrolled 8 times
L-1 has 1 loads, 1 stores, 4 prefetches, 0 FPadds,
1 FPmuls, and 0 FPdivs per iteration

UltraSPARC

A closer look ...

- ❑ on x86_64 Linux, we have no compiler that does give us the information we want ... i.e. the number of fpmul and fpdiv
- ❑ What now???
- ❑ We need some other tools, to get this information.
 - ❑ objdump – a Linux tool “to look into” executables
 - ❑ objdump has a lot of options – we need just one
 - ❑ `objdump -S file.o` : show disassembly and source intermixed
 - ❑ ... but how to interpret the assembly code???

A closer look ...

The really quick guide to x86 assembly

- ❑ instructions of interest: double precision floating point multiplication and division
- ❑ multiplication: `*mul*d`
 - ❑ `mulsd, mulpd, vmulsd, vmulpd`
- ❑ division: `*div*d`
 - ❑ `divsd, divpd, vdivsd, vdivpd`
- ❑ 's' is for single data, 'p' for packed data (vector)
- ❑ the leading letter indicates vector type

A closer look ...

compiled with 'suncc -g -fast -fsimple=0':

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
    9d:      c5 e3 5e e0                vdivsd %xmm0,%xmm3,%xmm4
    b8:      c5 d3 5e f0                vdivsd %xmm0,%xmm5,%xmm6
    c8:      c5 43 5e c0                vdivsd %xmm0,%xmm7,%xmm8
    d8:      c5 33 5e d0                vdivsd %xmm0,%xmm9,%xmm10
    e8:      c5 23 5e e0                vdivsd %xmm0,%xmm11,%xmm12
    f8:      c5 13 5e f0                vdivsd %xmm0,%xmm13,%xmm14
   108:      c5 83 5e c8                vdivsd %xmm0,%xmm15,%xmm1
   118:      c5 eb 5e d8                vdivsd %xmm0,%xmm2,%xmm3
   149:      c5 f3 5e d0                vdivsd %xmm0,%xmm1,%xmm2
   1cf:      c5 35 5e d6                vdivpd %ymm6,%ymm9,%ymm10
   1ed:      c5 25 5e e6                vdivpd %ymm6,%ymm11,%ymm12
   21a:      c5 45 5e c6                vdivpd %ymm6,%ymm7,%ymm8
   259:      c5 db 5e e8                vdivsd %xmm0,%xmm4,%xmm5
```

❑ only divisions!

❑ **single data** and **packed**: multi-versioning of loop

A closer look ...

compiled with 'suncc -g -fast -fsimple=2':

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
82:      c5 d3 5e f0          vdivsd %xmm0,%xmm5,%xmm6
a8:      c5 4b 59 04 06      vmulsd (%rsi,%rax,1),%xmm6,%xmm8
ba:      c5 4b 59 4c 06 08    vmulsd 0x8(%rsi,%rax,1),%xmm6,%xmm9
ca:      c5 4b 59 54 06 10    vmulsd 0x10(%rsi,%rax,1),%xmm6,%xmm10
d6:      c5 4b 59 5c 06 18    vmulsd 0x18(%rsi,%rax,1),%xmm6,%xmm11
e2:      c5 4b 59 64 06 20    vmulsd 0x20(%rsi,%rax,1),%xmm6,%xmm12
ee:      c5 4b 59 6c 06 28    vmulsd 0x28(%rsi,%rax,1),%xmm6,%xmm13
fa:      c5 4b 59 74 06 30    vmulsd 0x30(%rsi,%rax,1),%xmm6,%xmm14
106:     c5 4b 59 7c 06 38    vmulsd 0x38(%rsi,%rax,1),%xmm6,%xmm15
130:     c5 cb 59 3c 06      vmulsd (%rsi,%rax,1),%xmm6,%xmm7
19b:     c4 41 45 5e c8      vdivpd %ymm8,%ymm7,%ymm9
1cf:     c4 41 1d 59 e9      vmulpd %ymm9,%ymm12,%ymm13
1ee:     c4 41 0d 59 f9      vmulpd %ymm9,%ymm14,%ymm15
21a:     c4 41 2d 59 d9      vmulpd %ymm9,%ymm10,%ymm11
25e:     c5 e3 5e c0        vdivsd %xmm0,%xmm3,%xmm0
270:     c5 fb 59 24 06      vmulsd (%rsi,%rax,1),%xmm0,%xmm4
```

❑ mostly multiplications – it works!

A closer look ...

compiled with 'suncc -g -fast -fsimple=2'

❑ ... and using the 'restrict' keyword:

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
7d:      c4 41 35 5e da      vdivpd %ymm10,%ymm9,%ymm11
af:      c4 41 0d 59 fb      vmulpd %ymm11,%ymm14,%ymm15
ce:      c4 c1 75 59 d3      vmulpd %ymm11,%ymm1,%ymm2
fa:      c4 41 1d 59 eb      vmulpd %ymm11,%ymm12,%ymm13
140:     c5 db 5e e8      vdivsd %xmm0,%xmm4,%xmm5
150:     c4 a1 53 59 34 0e    vmulsd (%rsi,%r9,1),%xmm5,%xmm6
```

- ❑ mostly multiplications – it works!
- ❑ the multi-versioning is gone – only the vectorized code (and a clean-up loop) left!

A closer look ... now with GCC

- ❑ -O3 does not enable math optimizations, like replacing a division with a constant by a multiplication with the inverse
- ❑ to enable this, we need -ffast-math option
- ❑ to get some optimization information (at compile time), we can use -fopt-info

A closer look ... now with GCC

compiled with 'gcc -g -O3 -fopt-info'

```
$ gcc -g -O3 -fopt-info -c divvec.c
divvec.c:11:5: optimized: loop vectorized using 16 byte
vectors
divvec.c:11:5: optimized: loop versioned for vectorization
because of possible aliasing
```

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
35:      66 0f 5e c8          divpd   %xmm0,%xmm1
55:      f2 0f 5e c2          divsd   %xmm2,%xmm0
70:      f2 0f 5e c2          divsd   %xmm2,%xmm0
```

- ❑ GCC does multi-versioning, too!
- ❑ no multiplications – only divisions!

A closer look ... now with GCC

compiled with 'gcc -g -O3 -ffast-math -fopt-info'

```
$ gcc -g -O3 -ffast-math -fopt-info -c divvec.c
divvec.c:11:5: optimized: loop vectorized using 16 byte
vectors
divvec.c:11:5: optimized: loop versioned for vectorization
because of possible aliasing
```

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
1e:      f2 0f 5e c1          divsd    %xmm1,%xmm0
45:      66 0f 59 ca          mulpd    %xmm2,%xmm1
60:      f2 0f 59 04 c6      mulsd    (%rsi,%rax,8),%xmm0
80:      f2 0f 59 c8          mulsd    %xmm0,%xmm1
```

- ❑ GCC does multi-versioning, too!
- ❑ mostly multiplications – it works!

A closer look ... now with GCC

compiled with 'gcc -g -O3 -ffast-math -fopt-info'

❑ ... and using the 'restrict' keyword

```
$ gcc -g -O3 -ffast-math -fopt-info -c divvec.c
divvec.c:11:5: optimized: loop vectorized using 16 byte
vectors
```

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
10:      f2 0f 5e c1      divsd    %xmm1,%xmm0
35:      66 0f 59 ca      mulpd    %xmm2,%xmm1
52:      f2 0f 59 04 c6    mulsd    (%rsi,%rax,8),%xmm0
```

❑ multi-versioning is gone!

❑ less instructions than with the suncc compiler!?!

A closer look ... now with GCC

GCC does not unroll loops with -O3!

- ❑ we need to add -funroll-loops, too!

```
gcc -g -O3 -ffast-math -funroll-loops -fopt-info -c divvec.c
divvec.c:11:5: optimized: loop vectorized using 16 byte
vectors
divvec.c:11:16: optimized: loop unrolled 7 times
```

```
$ objdump -S divvec.o | egrep "mul.d|div.d" | grep -c mulpd
15
```

- ❑ now we have more mulpd instructions (14 more from loop unrolling)
- ❑ Lesson learned: check what '-O3' does!
 - ❑ Reminder: 'gcc -g -Q --help=optimizers ...'

A closer look ... summary

- ❑ Even with no compiler commentary, we can still get useful information, using
 - ❑ extra compiler options like -fopt-info (gcc)
 - ❑ tools like objdump – plus some basic knowledge about assembly code
- ❑ Caveat: this kind of analysis is feasible on small code kernels, only!
- ❑ Best practice: extract a small kernel from larger application, do the tests/tuning (- and reinsert).

Summary

- ❑ You have now heard about
 - ❑ tuning techniques
 - ❑ tools: compilers, analysis tools
 - ❑ libraries
 - ❑ other performance parameters
 - ❑ debuggers: try Totalview
- ❑ Now you have to apply that and get experience!
- ❑ But never forget:

Correct code has the highest
priority – not speed!