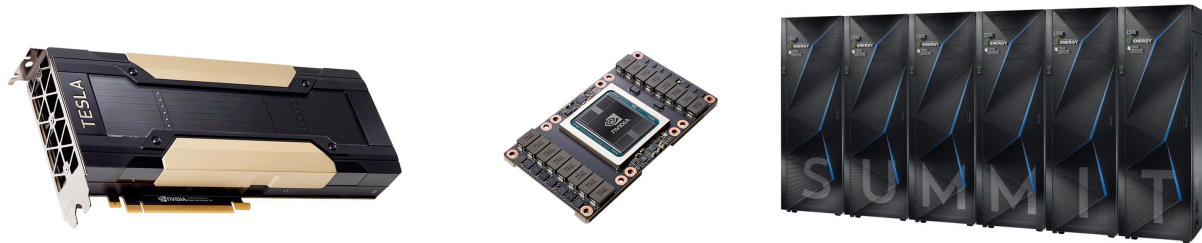


GPU Programming with OpenMP

Part 2: Data mapping

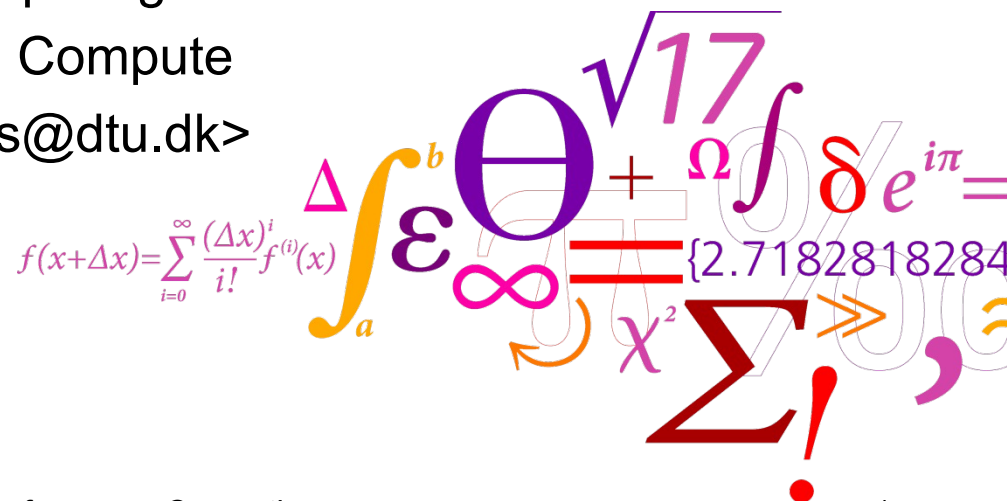


Hans Henrik Brandenburg Sørensen

DTU Computing Center

DTU Compute

<hhbs@dtu.dk>

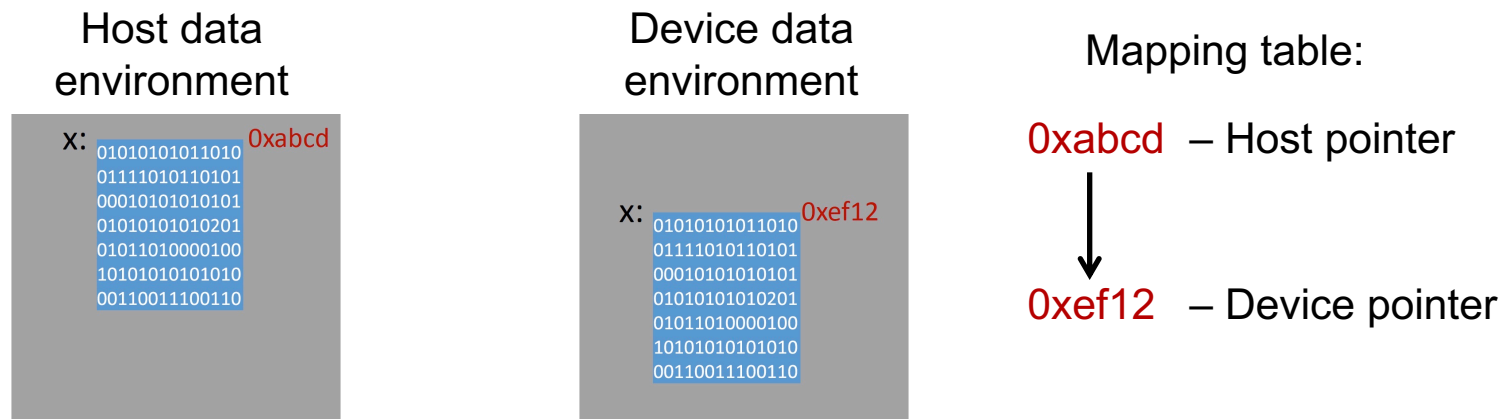


Overview

- OpenMP data mapping
 - `map` clause
 - `target` data
 - `target enter / exit` data
 - `target update`
 - `declare target`
- OpenMP array sections
- OpenMP runtime library
 - Device versions of `malloc`, `free`, `memcpy`
 - Matrix allocation on the device

OpenMP data environment

- OpenMP maintains a device data environment and a mapping table that records what memory pointers have been mapped from the host



- The table also maintains the translation between host memory pointers and device memory pointers
- Presence checks determine if data is already part of the device data environment before a transfer is made

OpenMP offload basics

■ target clauses related to data environment

- ☐ map
- ☐ is_device_ptr
- ☐ has_device_addr
- ☐ defaultmap
- ☐ ~~allocate~~
- ☐ ~~uses_allocators*~~

Not currently
supported in `nvc++`

OpenMP offload syntax

■ Syntax C/C++:

```
map([ [map-type-modifier[,] [map-type-  
modifier[,] ...] map-type : ] list)
```

- The `map` clause specifies how a list item is mapped from the host data environment to the corresponding list item in the device data environment

OpenMP offload syntax

- `map-type` can be

- `to`

- Allocates memory and moves data to the device

- `from`

- Allocates memory and moves data from the device

- `tofrom`

- Allocates memory and moves data to and from the device

- `alloc`

- Allocates memory on the device

- `release`

- The reference count is decreased by one

- `delete`

- Deletes data from the device (the reference count is set to 0)

OpenMP offload syntax

■ `map-type-modifier` can be

□ `always`

- Always copies the data to and from the device

□ ~~`close`~~

- A hint to the runtime to allocate memory close to the target device

□ ~~`mapper(mapper-identifier)`~~

- Use a user-defined mapper

□ ~~`present`~~

- This clause is ordered before other map clauses w/o present

□ ~~`iterator(iterator-definition)`~~

- Reference iterators defined by an iterators-definition

Not currently
supported in `nvc++`

OpenMP offload syntax

■ Example

```
double a[N], b[N], c[N];  
/*  
    initialize arrays  
*/  
#pragma omp target teams loop map(to: a, b) map(from: c)  
for (int i = 0; i < N; i++)  
    c[i] = a[i] + b[i];
```

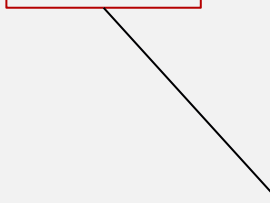
- ❑ Arrays a and b are transferred **to** the device before the offload
- ❑ Array c is **allocated** on the device before the offload
- ❑ Array c is transferred **from** the device after the offload

OpenMP offload syntax

■ Example – dynamic allocation

```
double *a, *b, c*;
/*
    allocate and initialize arrays of size N
*/
#pragma omp target teams loop \
    map(to: a[0:N], b[0:N]) map(from: c[0:N])
for (int i = 0; i < N; i++)
    c[i] = a[i] + b[i];

/* free arrays */
```



OpenMP array sections

- ❑ Arrays a and b are transferred **to** the device before the offload
- ❑ Array c is **allocated** on the device before the offload
- ❑ Array c is transferred **from** the device after the offload

OpenMP array sections

- An array section designates a subset of the elements in an array given by:
`[lower-bound : length : stride]`
 - ❑ Must be a subset of the original array
 - ❑ Array sections are allowed on multidimensional arrays
- Must be integers or integer expressions
 - ❑ If `lower-bound` is left out it defaults to 0
 - ❑ `length` must evaluate to a non-negative integer and must be explicitly specified when the size of the array dimension is not known
 - ❑ `stride` must evaluate to a positive integer (default 1)

OpenMP array sections

■ Example – dynamic allocation

```
double *a, *b, c*;
/*
    allocate and initialize arrays of size N
*/
#pragma omp target teams loop \
    map(to: a[0:N/2], b[0:N/2]) map(from: c[0:N/2])
for (int i = 0; i < N/2; i++)
    c[i] = a[i] + b[i];

#pragma omp target teams loop \
    map(to: a[N/2:N/2], b[N/2:N/2]) map(from: c[N/2:N/2])
for (int i = N/2; i < N; i++)
    c[i] = a[i] + b[i];
```

OpenMP offload syntax

■ Another example

```
double a[N], b[N], c[N], d[N];
/*
    initialize arrays
*/
#pragma omp target teams loop map(to: a, b) map(from: c)
for (int i = 0; i < N; i++)
    c[i] = a[i] + b[i];

#pragma omp target teams loop map(to: a, b) map(from: d)
for (int i = 0; i < N; i++)
    d[i] = a[i] - b[i];
```

- ❑ Two offload regions that both use a and b (transferred 2 times)
- ❑ This works, but it is not the best we can do!

OpenMP offload syntax

■ Syntax C/C++:

```
#pragma omp target data [clause]
{
    ...
}
```

■ Clause can be

- ❑ `if([target data :] scalar_expr)`
- ❑ `device(int_expr)`
- ❑ `map(...)`
- ❑ `use_device_ptr(ptr-list)`
- ❑ `use_device_addr(list)`

■ Creates persistent data environment within { }

OpenMP offload syntax

■ So the better solution is

```
double a[N], b[N], c[N], d[N];
/*
    initialize arrays
*/
#pragma omp target data map(to: a, b) map(from: c, d)
{
    #pragma omp target teams loop map(to: a, b) map(from: c)
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];

    #pragma omp target teams loop map(to: a, b) map(from: d)
    for (int i = 0; i < N; i++)
        d[i] = a[i] - b[i];
}
```

Presence check: a and b already available!

□ Arrays a and b are transferred once and used in both offloads regions

OpenMP measuring runtimes

■ Measuring the transfer time

```
double a[N], b[N], c[N];
/*
    initialize arrays and warmup device
*/
double t = omp_get_wtime();
#pragma omp target data map(to: a, b) map(from: c)
{
    double t = omp_get_wtime();
    #pragma omp target teams loop map(to: a, b) map(from: c)
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];
    printf("Runtime w/o transfer: %f\n", omp_get_wtime() - t);
}
printf("Runtime with transfer: %f\n", omp_get_wtime() - t);
```

```
$ ./vecadd
Runtime w/o transfer: 0.000565
Runtime with transfer: 0.002533
```

❑ For reliable runtimes run offload region several times and average

OpenMP offload syntax

■ Syntax C/C++:

```
#pragma omp target enter data [clause]  
#pragma omp target exit data [clause]
```

■ Clause can be

- ❑ `if([target data :] scalar_expr)`
- ❑ `device(int_expr)`
- ❑ `map(...)`
- ❑ `depend([depend-modifier,] depen-type : list)`
- ❑ `nowait`

■ Standalone directives that specifies mapping to the data environment of the default device

OpenMP offload syntax

■ Target enter / exit example

main.cpp

```
int main(int argc, char *argv[]) {
    double a[N], b[N], c[N];

    initialize(a, b, c);

    #pragma omp target teams loop \
        map(to: a, b) map(from: c)
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];

    finalize(a, b, c);

    return(0);
}
```

functions.cpp

```
void initialize(double *a, double *b, double *c) {
    for (int i = 0; i < N; i++)
        a[i] = b[i] = i;
    #pragma omp target enter data \
        map(to: a[:N], b[:N]) map(alloc: c[:N])
}

void finalize(double *a, double *b, double *c) {
    #pragma omp target exit data \
        map(release: a[:N], b[:N]) map(from: c[:N])
    for (int i = 0; i < N; i++)
        printf("%f\n", c[i]);
}
```

□ Remember to release a and b in order to clean up data environment

OpenMP measuring runtime

■ H2D and D2H transfers

main.cpp

```
int main(int argc, char *argv[]) {
    double a[N], b[N], c[N];
    /*
     * warm up device
     */
    initialize(a, b, c);

    #pragma omp target teams loop \
        map(to: a, b) map(from: c)
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];

    finalize(a, b, c);

    return(0);
}
```

functions.cpp

```
void initialize(double *a, double *b, double *c) {
    for (int i = 0; i < N; i++)
        a[i] = b[i] = i;
    double t = omp_get_wtime();
    #pragma omp target enter data \
        map(to: a[:N], b[:N]) map(alloc: c[:N])
    printf("H2D transfer time: %f\n",
        omp_get_wtime() - t);
}

void finalize(double *a, double *b, double *c) {
    double t = omp_get_wtime();
    #pragma omp target exit data \
        map(release: a[:N], b[:N]) map(from: c[:N])
    printf("D2H transfer time: %f\n",
        omp_get_wtime() - t);
}
```

```
$ ./vecadd
H2D transfer time: 0.001761
D2H transfer time: 0.001139
```

❑ Remember to warm up the device or transfer times will be +0.2 secs

OpenMP offload syntax

■ Syntax C/C++:

```
#pragma omp target data update [clause]
```

■ Clause can be

- ❑ `if([target data :] scalar_expr)`
- ❑ `device(int_expr)`
- ❑ `depend([depend-modifier,] depen-type : list)`
- ❑ `nowait`
- ❑ `to([motion-modifier[,] ...]:] list)`
- ❑ `from([motion-modifier[,] ...]:] list)`

■ Standalone directive that makes the device data environment consistent with their original list items, according to the specified motion clauses

OpenMP offload syntax

■ Target update example

```
#pragma omp target data map(to: a, b) map(from: c, d)
{
    #pragma omp target teams loop map(to: a, b) map(from: c)
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];

    modify_b_on_host(b);
    #pragma omp target update to(b)

    #pragma omp target teams loop map(to: a, b) map(from: d)
    for (int i = 0; i < N; i++)
        d[i] = a[i] - b[i];
}
```

- Array b is modified in the host data environment and needs to be updated in the device data environment before the final offload

OpenMP offload syntax

■ Syntax C/C++:

```
#pragma omp declare target  
declarations  
#pragma omp end declare target
```

```
#pragma omp declare target [(list) | clause]
```

■ Clause can be

- ❑ `to(list)`
- ❑ ~~`link(list)`~~
- ❑ `device_type(host | nohost | any)`

- ## ■ This directive specifies that global variables and functions (C/C++) are mapped to a device for all device executions (or for a specific one via link)

OpenMP offload syntax

■ Using declare target

```
#pragma omp declare target
double a[N], b[N], c[N];
#pragma omp end declare target

void vecadd() {
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];
}

#pragma omp declare target (vecadd)

int main(int argc, char *argv[]) {
    /* initialize arrays on host */
    vecadd(); // Call host version
    #pragma omp target update to(a, b)
    #pragma omp target
    vecadd(); // Call device version
    #pragma omp target update from(c)
}
```

← no mapping – arrays are globally available

Implicit data mapping

■ Why did this run?

```
#define N 16
int main(int argc, char *argv[]) {
    double a[N], b[N], c[N];
    for (int i = 0; i < N; i++)
        a[i] = b[i] = i * 1.0;

    #pragma omp target teams \
        distribute parallel for
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];

    for (int i = 0; i < N; i++)
        printf("%f\n", c[i]);
}
```

← arrays are implicitly mapped to the device

← arrays are implicitly mapped back to the host

Implicit data mapping

- If a variable is a scalar then it is implicitly `firstprivate` (i.e., not mapped)
- If a variable is not a scalar then implicitly a `map (tofrom: ...)` is added
- If a variable is a pointer it “is treated as if it is the base pointer of a zero-length array section that had appeared as a list item in a `map` clause”
- However, if the `defaultmap` clause is present in the construct, it takes precedence

Summary of data mapping

- Map clause on a target construct
 - Maps variables for a single target region
 - Enclosed region executes on device and maps data
- Target data
 - Map variables across multiple target regions
 - Region does not execute on device, only maps data
- Declare target
 - Mapping variables and functions for the whole execution of the program (“globally mapped”)
- Target enter/exit/update
 - Map variables in stand-alone clauses

OpenMP runtime library

OpenMP offload runtime library

■ Managing memory yourself

Name

```
void* omp_target_alloc(size_t,  
int dev_num)
```

```
void omp_target_free(void*,  
int dev_num)
```

```
int omp_target_memcpy(...,  
int dev_num)
```

```
int omp_target_memcpy_rect(...,  
int dev_num)
```

```
omp_target_associate_ptr(...)
```

```
omp_target_disassociate_ptr(...)
```

```
omp_target_is_present(...)
```

Functionality

allocate memory on device

free memory on device

memcpy to and from device

memcpy to and from device of
a rectangular subvolume

combining device ptr with

host ptr to be used in map

clause

(we use is_device_ptr clause)

OpenMP offload syntax

■ Syntax C/C++:

```
void* omp_target_alloc(size_t size,  
                        int dev_num);
```

- This routine allocates memory of `size` bytes in the device data environment of device `dev_num` and returns a device pointer to that memory

■ Syntax C/C++:

```
void omp_target_free(void *dev_ptr,  
                     int dev_num);
```

- This routine frees the memory at `dev_ptr` in the device data environment of device `dev_num`

OpenMP offload syntax

■ Syntax C/C++:

```
int omp_target_memcpy(void *dst,  
                      void *src,  
                      size_t length,  
                      size_t dst_offset,  
                      size_t src_offset,  
                      int dst_dev_num,  
                      int src_dev_num);
```

- This routine copies `length` bytes of memory at offset `src_offset` from `src` in the device data environment of device `src_dev_num` to `dst` starting at offset `dst_offset` in the device data environment of device `dst_dev_num`

OpenMP offload syntax

- The `is_device_ptr` clause indicates that its list items are device pointers
 - Inside the target construct, each list item is privatized and the new list item is initialized to the device address to which the original list item refers
- Support for device pointers created outside of OpenMP, specifically outside of any OpenMP mechanism that returns a device pointer, is implementation defined

CUDA device pointers are fully supported in `nvc++`

OpenMP offload syntax

■ Example – dynamic allocation

```
double *a, *b, *c, *a_d, *b_d, *c_d; // Notation _d is convention for device pointers.
/* allocate and initialize arrays on host */

int dev_num = omp_get_default_device();
a_d = (double*)omp_target_alloc(N * sizeof(double), dev_num);
b_d = (double*)omp_target_alloc(N * sizeof(double), dev_num);
c_d = (double*)omp_target_alloc(N * sizeof(double), dev_num);

omp_target_memcpy(a_d, a, N * sizeof(double), 0, 0, dev_num, omp_get_initial_device());
omp_target_memcpy(b_d, b, N * sizeof(double), 0, 0, dev_num, omp_get_initial_device());

#pragma omp target teams loop is_device_ptr(a_d, b_d, c_d)
for (int i = 0; i < N; i++)
    c_d[i] = a_d[i] + b_d[i];

omp_target_memcpy(c, c_d, N * sizeof(double), 0, 0, omp_get_initial_device(), dev_num);
...
omp_target_free(a_d, dev_num);
omp_target_free(b_d, dev_num);
omp_target_free(c_d, dev_num);
```

clause `is_device_ptr` required here

■ Example matrix allocation – host version

```
/* Routine for allocating two-dimensional array */
double **malloc_2d(int m, int n) {
    if (m <= 0 || n <= 0)
        return NULL;

    double **A = (double**)malloc(m * sizeof(double *));
    if (A == NULL)
        return NULL;

    A[0] = (double*)malloc(m * n * sizeof(double));
    if (A[0] == NULL) {
        free(A);
        return NULL;
    }

    for (int i = 0; i < m; i++)
        A[i] = A[0] + i * n;

    return A;
}
```

```
void free_2d(double **A) {
    free(A[0]);
    free(A);
}
```


■ Example matrix allocation – device version

```
/* Routine for allocating two-dimensional array on the device */
double **malloc_2d_dev(int m, int n, double **data) {
    if (m <= 0 || n <= 0)
        return NULL;

    double **A = (double**)omp_target_alloc(m*sizeof(double *), omp_get_default_device());
    if (A == NULL)
        return NULL;

    double *a = (double*)omp_target_alloc(m*n * sizeof(double), omp_get_default_device());
    if (a == NULL) {
        omp_target_free(A, omp_get_default_device());
        return NULL;
    }
    #pragma omp target is_device_ptr(A, a)
    for (int i = 0; i < m; i++)
        A[i] = a + i * n;

    *data = a; return A;
}
```

```
void free_2d_dev(double **A,
                 double *data) {
    omp_target_free(data,
                    omp_get_default_device());
    omp_target_free(A,
                    omp_get_default_device());
}
```

■ Example matrix allocation

```
double **A = malloc_2d(N, N); // Allocate A on host

double *data;
double **A_d = malloc_2d_dev(N, N, &data); // Allocate A_d on device

/* initialize A on host */

omp_target_memcpy(data, A[0], N * N * sizeof(double), // Copy data from A to A_d
                  0, 0, omp_get_default_device(), omp_get_initial_device());

double sum = 0.0;
#pragma omp target teams distribute parallel for reduction(+:sum) is_device_ptr(A_d)
    for (int i = 0; i < N; i++)
        for (int j = 0; j < N; j++)
            sum += A_d[i][j] * A_d[i][j];
printf("Frobenius norm: %f\n", sqrt(sum));

free_2d(A); // Free A on host
free_2d_dev(A_d, data); // Free A_d on device
```

Exercises

- Finish up the first two exercises
 - ex1_deviceQuery
 - ex2_helloworld
- Begin the third exercise
 - ex3_mandelbrot

End of lecture